

1.

```
[hive> create table payments (VendorID string, tprep_pickup_datetime date, payment_type string, total_amount float ) row format delimited fields terminated by ',' stored as textfile;
OK
Time taken: 1.56 seconds
```

```
[hive> show tables;
OK
congestion
distance
passengers
payments
tolls
tripdata_table
Time taken: 0.127 seconds, Fetched: 6 row(s)
hive> █
```

2.

```
[hive> describe passengers;
OK
tprep_pickup_datetime  date
passenger_count        int
total_amount           float
Time taken: 0.322 seconds, Fetched: 3 row(s)
hive> }█
```

```
OK
tprep_pickup_datetime  date
passenger_count        int
trip_distance          float
total_amount           float
Time taken: 0.194 seconds, Fetched: 4 row(s)
hive> █
```

3.

```
hadoop@34d13db60376:/home$ cd hadoop
hadoop@34d13db60376:~$ ls
airflow      hadoopdata      hs_err_pid10887.log  nohup.out      spark-warehouse
codegen_region.java  hive            hs_err_pid11254.log  region.java     sqoop
derby.log        hs_err_pid10630.log  landing             scripts         yarn-utils.py
hadoop          hs_err_pid10724.log  metastore_db        spark
hadoop@34d13db60376:~$ wget https://dataengineerpublic.blob.core.windows.net/data-engineer/yellow_tripdata_2021-01.csv -O yellow_tripdata_2021-01.csv
[--2024-08-17 18:30:54-- https://dataengineerpublic.blob.core.windows.net/data-engineer/yellow_tripdata_2021-01.csv
Resolving dataengineerpublic.blob.core.windows.net (dataengineerpublic.blob.core.windows.net)... 20.150.25.164
Connecting to dataengineerpublic.blob.core.windows.net (dataengineerpublic.blob.core.windows.net)|20.150.25.164|:443... connected
.
HTTP request sent, awaiting response... 200 OK
Length: 125981363 (120M) [text/csv]
Saving to: 'yellow_tripdata_2021-01.csv'

yellow_tripdata_2021-01. 100%[=====] 120.14M  2.90MB/s   in 26s

2024-08-17 18:31:21 (4.65 MB/s) - 'yellow_tripdata_2021-01.csv' saved [125981363/125981363]

hadoop@34d13db60376:~$ ls
airflow      hive            landing         spark
codegen_region.java  hs_err_pid10630.log  metastore_db    spark-warehouse
derby.log      hs_err_pid10724.log  nohup.out       sqoop
hadoop          hs_err_pid10887.log  region.java     yarn-utils.py
hadoopdata     hs_err_pid11254.log  scripts         yellow_tripdata_2021-01.csv
hadoop@34d13db60376:~$ hdfs dfs -ls /
Found 8 items
drwxrwxrwx - hadoop supergroup      0 2024-08-15 09:33 /ingest
drwxr-xr-x - hadoop supergroup      0 2022-04-26 19:51 /inputs
drwxr-xr-x - hadoop supergroup      0 2022-01-22 21:35 /logs
drwxrwxrwx - hadoop supergroup      0 2024-08-11 22:36 /nifi
drwxr-xr-x - hadoop supergroup      0 2024-08-11 20:30 /sqoop
drwxr-xr-x - hadoop supergroup      0 2024-08-09 10:33 /table
drwxrwxr-x - hadoop supergroup      0 2022-05-02 20:46 /tmp
drwxr-xr-x - hadoop supergroup      0 2022-01-23 13:15 /user
hadoop@34d13db60376:~$ hdfs dfs -put yellow_tripdata_2021-01.csv /ingest
hadoop@34d13db60376:~$ hdfs dfs -ls /ingest
Found 1 items
-rw-r--r-- 1 hadoop supergroup 125981363 2024-08-17 18:33 /ingest/yellow_tripdata_2021-01.csv
```



version 3.2.0

```
Using Python version 3.8.10 (default, Mar 15 2022 12:22:08)
Spark context Web UI available at http://34d13db60376:4040
Spark context available as 'sc' (master = yarn, app id = application_1723931863203_0001).
SparkSession available as 'spark'.
>>> df = spark.read.csv("hdfs://172.17.0.2:9000/ingest/yellow_tripdata_2021-01.csv", header=True, inferSchema=True)
>>> df.Schema()
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
    raise AttributeError(
AttributeError: 'DataFrame' object has no attribute 'Schema'
>>> df.printSchema()
root
 |-- VendorID: integer (nullable = true)
 |-- tpep_pickup_datetime: string (nullable = true)
 |-- tpep_dropoff_datetime: string (nullable = true)
 |-- passenger_count: integer (nullable = true)
 |-- trip_distance: double (nullable = true)
 |-- RatecodeID: integer (nullable = true)
 |-- store_and_fwd_flag: string (nullable = true)
 |-- PULocationID: integer (nullable = true)
 |-- DOLocationID: integer (nullable = true)
 |-- payment_type: integer (nullable = true)
 |-- fare_amount: double (nullable = true)
 |-- extra: double (nullable = true)
 |-- mta_tax: double (nullable = true)
 |-- tip_amount: double (nullable = true)
 |-- tolls_amount: double (nullable = true)
 |-- improvement_surcharge: double (nullable = true)
 |-- total_amount: double (nullable = true)
 |-- congestion_surcharge: double (nullable = true)
```

5.

```
>>> df_payments= df.select("VendorId", "tpep_pickup_datetime", "payment_type", "total_amount")
>>> df_payments.show()
```

VendorId	tpep_pickup_datetime	payment_type	total_amount
1	2021-01-01 00:30:10	2	11.8
1	2021-01-01 00:51:20	2	4.3
1	2021-01-01 00:43:30	1	51.95
1	2021-01-01 00:15:48	1	36.35
2	2021-01-01 00:31:49	1	24.36
1	2021-01-01 00:16:29	1	14.15
1	2021-01-01 00:00:28	2	17.3
1	2021-01-01 00:12:29	2	21.8
1	2021-01-01 00:39:16	4	28.8
1	2021-01-01 00:26:12	1	18.95
2	2021-01-01 00:15:52	1	24.3
2	2021-01-01 00:46:36	1	10.79
1	2021-01-01 00:10:46	2	33.92
2	2021-01-01 00:31:06	1	14.16
2	2021-01-01 00:42:11	2	8.3
2	2021-01-01 00:17:48	1	10.3
2	2021-01-01 00:33:38	1	12.09
2	2021-01-01 00:47:56	1	12.36
2	2021-01-01 00:04:21	1	9.96
2	2021-01-01 00:18:36	2	12.3

only showing top 20 rows

```
NameError: name payment_type is not defined
```

```
>>> df_payments_tcredito= df_payments.filter(df_payments["payment_type"]==1)
>>> df_payments_tcredito.show()
```

VendorId	tpep_pickup_datetime	payment_type	total_amount
1	2021-01-01 00:43:30	1	51.95
1	2021-01-01 00:15:48	1	36.35
2	2021-01-01 00:31:49	1	24.36
1	2021-01-01 00:16:29	1	14.15
1	2021-01-01 00:26:12	1	18.95
2	2021-01-01 00:15:52	1	24.3
2	2021-01-01 00:46:36	1	10.79
2	2021-01-01 00:31:06	1	14.16
2	2021-01-01 00:17:48	1	10.3
2	2021-01-01 00:33:38	1	12.09
2	2021-01-01 00:47:56	1	12.36
2	2021-01-01 00:04:21	1	9.96
2	2021-01-01 00:56:30	1	11.84
1	2021-01-01 00:37:59	1	30.8
2	2021-01-01 00:34:37	1	18.3
2	2021-01-01 00:06:24	1	22.8
2	2021-01-01 00:35:17	1	26.16
2	2021-01-01 00:13:44	1	22.88
2	2021-01-01 00:43:03	1	11.0
2	2021-01-01 00:19:57	1	40.3

only showing top 20 rows

```
>>> df_payments_tcredito.write.insertInto("tripdata.payments")
2024-08-17 20:00:29,192 WARN conf.HiveConf: HiveConf of name hive.metastore.local does not exist
2024-08-17 20:00:31,303 WARN session.SessionState: METASTORE_FILTER_HOOK will be ignored, since hive.security.authorization.manager is set to instance of HiveAuthorizerFactory.
>>>
```



```

-----
Time taken: 0.188 seconds, Fetched: 6 row(s)
hive> select * from payments limit 10;
OK
1      2021-01-01      1      51.95
1      2021-01-01      1      36.35
2      2021-01-01      1      24.36
1      2021-01-01      1      14.15
1      2021-01-01      1      18.95
2      2021-01-01      1      24.3
2      2021-01-01      1      10.79
2      2021-01-01      1      14.16
2      2021-01-01      1      10.3
2      2021-01-01      1      12.09
Time taken: 2.215 seconds, Fetched: 10 row(s)
hive> █

```

6.

```

# Hive security, authorization, manager is set to instance of HiveAuthorizationManager.
[>>> df_passengers= df.select(to_date(col("tpep_pickup_datetime"), "yyyy-MM-dd HH:mm:ss"),col("passenger_count").cast("int"),col("total_amount").cast("float"))
[>>> df_passengers.printSchema()
root
|-- to_date(tpep_pickup_datetime, yyyy-MM-dd HH:mm:ss): date (nullable = true)
|-- passenger_count: integer (nullable = true)
|-- total_amount: float (nullable = true)

[>>> █

[>>> df_passengers_filt = df_passengers.filter(
...     (df_passengers["passenger_count"] > 2) & (df_passengers["total_amount"] > 8)
... )
[>>> df_passengers_filt.show()
+-----+-----+-----+
|to_date(tpep_pickup_datetime, yyyy-MM-dd HH:mm:ss)|passenger_count|total_amount|
+-----+-----+-----+
|2021-01-01|3|24.3|
|2021-01-01|5|14.16|
|2021-01-01|5|8.3|
|2021-01-01|3|9.3|
|2021-01-01|4|18.3|
|2021-01-01|4|13.3|
|2021-01-01|3|40.3|
|2021-01-01|5|14.8|
|2021-01-01|3|18.59|
|2021-01-01|3|13.56|
|2021-01-01|3|9.96|
|2021-01-01|3|66.36|
|2021-01-01|3|15.95|
|2021-01-01|3|15.8|
|2021-01-01|3|13.3|
|2021-01-01|3|11.76|
|2021-01-01|3|31.8|
|2021-01-01|3|12.95|
|2021-01-01|3|10.8|
|2021-01-01|4|22.8|
+-----+-----+-----+
only showing top 20 rows

[>>> █

```

```
[>>> df_passengers_filt.write.insertInto("tripdata.passengers")
>>> █
```

```
[hive> select * from passengers limit 10;
```

```
OK
```

```
2021-01-01      3      24.3
2021-01-01      5      14.16
2021-01-01      5       8.3
2021-01-01      3       9.3
2021-01-01      4      18.3
2021-01-01      4      13.3
2021-01-01      3      40.3
2021-01-01      5      14.8
2021-01-01      3      18.59
2021-01-01      3      13.56
```

```
Time taken: 1.356 seconds, Fetched: 10 row(s)
```

```
hive> █
```

7.

only showing top 20 rows

```
[>>> df_passengers_filt.write.insertInto("tripdata.passengers")
>>> df_tolls=df.select(to_date(col("tpep_pickup_datetime"), "yyyy-MM-dd HH:mm:ss").alias("tpep_pickup_datetime"),col("passenger_count").cast("int").alias("passenger_count"),col("tolls_amount").cast("float").alias("tolls_amount"),col("total_amount").cast("float").alias("total_amount"))
>>> df_tolls.printSchema()
```

```
root
 |-- tpep_pickup_datetime: date (nullable = true)
 |-- passenger_count: integer (nullable = true)
 |-- tolls_amount: float (nullable = true)
 |-- total_amount: float (nullable = true)
```

```
>>> █
```

```
>>> df_tolls_filt = df_tolls.filter(
...     (df_tolls["tolls_amount"] > 0.1) & (df_tolls["passenger_count"] > 1)
... )
```

```
>>> █
```

```
[>>> df_tolls_filt.show(5)
```

tpep_pickup_datetime	passenger_count	tolls_amount	total_amount
2021-01-01	2	6.12	33.92
2021-01-01	2	6.12	59.42
2021-01-01	2	6.12	35.92
2021-01-01	6	6.12	40.1
2021-01-01	3	6.12	54.0

only showing top 5 rows

```
>>> █
```

only showing top 5 rows

```
[>>> df_tolls_filt.write.insertInto("tripdata.tolls")
```

```
>>> █
```

```
[hive> select * from tolls limit 10;
OK
2021-01-01      2      6.12      33.92
2021-01-01      2      6.12      59.42
2021-01-01      2      6.12      35.92
2021-01-01      6      6.12      40.1
2021-01-01      3      6.12      54.0
2021-01-01      2      2.8       34.1
2021-01-01      4      6.12      61.42
2021-01-01      4      6.12      51.42
2021-01-01      2     11.75     12.05
2021-01-01      6      6.12     71.42
Time taken: 0.981 seconds, Fetched: 10 row(s)
hive> █
```

8.

```
>>> df_congestion=df.select(to_date(col("tpep_pickup_datetime"), "yyyy-MM-dd HH:mm:ss").alias("tpep_pickup_datetime"),col("passenger_count").cast("int").alias("passenger_count"),col("congestion_surcharge").cast("float").alias("congestion_surcharge"),col("total_amount").cast("float").alias("total_amount"))
```

```
[>>> df_congestion_filt= df_congestion.filter(to_date(col("tpep_pickup_datetime"), "yyyy-MM-dd") == "2021-01-18")
```

```
[>>> df_congestion_filt.show(10)
```

tpep_pickup_datetime	passenger_count	congestion_surcharge	total_amount
2021-01-18	1	2.5	10.8
2021-01-18	1	2.5	16.56
2021-01-18	1	0.0	10.3
2021-01-18	1	2.5	11.16
2021-01-18	1	2.5	11.3
2021-01-18	1	2.5	21.23
2021-01-18	1	2.5	12.96
2021-01-18	1	2.5	13.87
2021-01-18	1	2.5	14.8
2021-01-18	1	2.5	14.14

```
only showing top 10 rows
```

```
NameError: name tripdata is not defined
```

```
[>>> df_congestion_filt.write.insertInto("tripdata.congestion")
```

```
>>> █
```



```
[hive> select * from congestion limit 5;
OK
2021-01-18      1      2.5      10.8
2021-01-18      1      2.5      16.56
2021-01-18      1      0.0      10.3
2021-01-18      1      2.5      11.16
2021-01-18      1      2.5      11.3
Time taken: 0.675 seconds, Fetched: 5 row(s)
hive> █
```

9.

```
[>>> df_distance=df.select(to_date(col("tpep_pickup_datetime"), "yyyy-MM-dd HH:mm:ss").alias("tpep_
pickup_datetime"),col("passenger_count").cast("int").alias("passenger_count"),col("trip_distance")
.cast("float").alias("trip_distance"),col("total_amount").cast("float").alias("total_amount"))
[>>> df_distance.show(5)
+-----+-----+-----+-----+
|tpep_pickup_datetime|passenger_count|trip_distance|total_amount|
+-----+-----+-----+-----+
|2021-01-01|1|2.1|11.8|
|2021-01-01|1|0.2|4.3|
|2021-01-01|1|14.7|51.95|
|2021-01-01|0|10.6|36.35|
|2021-01-01|1|4.94|24.36|
+-----+-----+-----+-----+
only showing top 5 rows

[>>> df_distance_filt=df_distance.filter((to_date(col("tpep_pickup_datetime"), "yyyy-MM-dd") == "22|
0-12-31") & (df_distance["passenger_count"] == 1) & (df_distance["trip_distance"] > 15 ))
[>>> df_
df_congestion      df_distance_filt      df_payments      df_tolls_filt
df_congestion_filt  df_passengers      df_payments_tcredit  df_tolls_filt
df_distance         df_passengers_filt  df_tolls
[>>> df_
df_congestion      df_distance_filt      df_payments      df_tolls_filt
df_congestion_filt  df_passengers      df_payments_tcredit  df_tolls_filt
df_distance         df_passengers_filt  df_tolls
[>>> df_distance_filt.show(5)
+-----+-----+-----+-----+
|tpep_pickup_datetime|passenger_count|trip_distance|total_amount|
+-----+-----+-----+-----+

>>> █

[>>> df_distance_filt.write.insertInto("tripdata.distance")
>>> █
Time taken: 0.000 seconds, Fetched: 4 row(s),
[hive> select * from distance limit 5;
OK
Time taken: 0.792 seconds
hive> █
```